


# Fantastic Hacks and Where to Find Them

Adapted from Miriam Pena's "[BEAM Extreme](#)" talk @ ElixirConf 2019  
Slides available at [tylerayoung.com](http://tylerayoung.com)

# 👋 Hi, I'm Tyler

- Writing Elixir since 2019
- Built X-Plane massive multiplayer game server
- Worked at Felt on collaborative mapmaking
- Now starting [SleepEasy Website Monitor](#)






**Make it work, then make it  
beautiful, then if you really, really  
have to, make it fast.**

Joe Armstrong, *Erlang & OTP in Action*

# But first, profile.

- Perf is usually good enough by default
  - Don't thrash without a reason
  - Don't write horrifying hacks without a *really* good reason




# fprof |> eflame

# Phoenix LiveDashboard

[Home](#) [OS Data](#) [Metrics](#) [Request Logger](#) [Applications](#) [Processes](#) [Ports](#) [Sockets](#) [ETS](#) [Flame On](#)

Enable



flame on

# It's probably the database.

- Database queries are the #1 driver of performance in modern web apps
- Beware N + 1 queries


```
Enum.map(event_ids, fn id →  
  event = load_event(id)  
  # do stuff  
end)
```

# It's probably the database.

- Database queries are the #1 driver of performance in modern web apps
- Beware N + 1 queries
- Indexing (and multi-column indexing) can help
  - 😎 Read speed (WHERE, ORDER\_BY, JOIN ON)
  - 😢 Disk space
  - 😱 Insert speed?

# Elixir: Message passing by default

- 👍 No shared memory, no problem!
- 😎 Concurrency is harder to screw up
- 😎 Garbage collection is easy!

A close-up shot of Red Skull, a character from the Marvel Cinematic Universe. He is wearing his signature red helmet with a black visor and a black suit underneath. He is looking slightly to the left with a serious expression. In the background, a person's shoulder and part of their face are visible, suggesting a conversation.

But what is message passing  
if not `memcpy` persevering?

# Elixir: Message passing by default

- 👍 No shared memory, no problem!
- 😎 Concurrency is harder to screw up
- 😎 Garbage collection is easy!
- 😤 All those copies eat CPU time

# Elixir: Message passing by default\*

- 👍 No shared memory, no problem!
- 😎 Concurrency is harder to screw up
- 😎 Garbage collection is easy!
- 😤 All those copies eat CPU time

\* Except atoms, binaries over 64 bytes, literals, & code

# Elixir: Message passing by default\*

- 👍 No shared memory, no problem!
- 😎 Concurrency is harder to screw up
- 😎 Garbage collection is easy!
- 😤 All those copies eat CPU time

*Foreshadowing*



\* Except atoms, binaries over 64 bytes, literals, & code

# Don't copy more than you have to

- Be smart about what you send between processes
  - Consider not separating out a process at all?
  - “Big Data” is calling from the year 2010:  
Bring your computation to the data, rather than the data to the function
  - Look for instances where a data structure is getting copied due to immutability guarantees

```
20 -   {far_index, _, far_squared_dist} =
21 -     Enum.zip(0..(length(segment) - 1), segment)
22 -   |> Enum.drop(1)
23 -   |> Enum.drop(-1)
24 -   |> Enum.map(fn {i, p} -> {i, p, seg_dist(p, first, last)} end)
25 -   |> Enum.max_by(&element(&1, 2))
22 +   {_, far_value, far_index, far_squared_dist} =
23 +     Enum.reduce(middle, {1, nil, 1, 0}, fn element, {idx, max_val, max_idx, max_dist} =>
24 +       dist = seg_dist(element, first, last)
25 +
26 +       if dist >= max_dist do
27 +         {idx + 1, element, idx, dist}
28 +       else
29 +         {idx + 1, max_val, max_idx, max_dist}
30 +       end
31 +     end)
```

[https://github.com/pkinney/simplify\\_ex/pull/4/files](https://github.com/pkinney/simplify_ex/pull/4/files)



Call Kenny Loggins, because  
you're in the *danger zone*!

# Choose a better data type

- Lower overhead
  - Prefer iolist to concatenation
  - Prefer BIFs & pattern matching to function calls


# Choose a better data type

- Lower overhead
  - Prefer iolist to concatenation
  - Prefer BIFs & pattern matching to function calls
- Better memory locality
  - Prefer binaries to (char)lists
  - Prefer tuples/records to lists/maps/structs? (

# Aside: Memory locality




# Aside: Memory locality




# Aside: Memory locality




# Aside: Why does memory locality matter?




# Aside: Why does memory locality matter?



# Aside: Why does memory locality matter?



# Aside: Why does memory locality matter?



# Aside: Why does memory locality matter?



Source  
[tylerayoung.com](http://tylerayoung.com)

# Choose a better data type

- Lower overhead
  - Prefer iolist to concatenation
  - Prefer BIFs & pattern matching to function calls
- Better memory locality
  - Prefer binaries to (char)lists
  - Prefer tuples/records to lists/maps/structs? (

# Move it to compile time

- Macros can turn computation & data into code
  - Can't change for the life of the program
- 😎 Not copied between processes!
- ⚠ Compile time cost
- ⚠ Theoretical runtime cost due to code bloat

# Use ETS to share data across processes

- Shared, in-memory table
- : protected by default (only writable by the process that creates it)
- 😎 Concurrent reads (no process bottleneck)
- 😎 Faster than GenServer
- 😐 Still have to copy data in & out!
- ⚠️ Exciting new possibilities for race conditions

# Skip GenServer in favor of “raw” processes

- GenServer is designed for high *availability*, not high throughput
  - “Raw” processes (`spawn/send/receive`) avoid layers of indirection
- 😎 ~80% faster per call!
- ⚠️ Code is harder to read, less idiomatic



# Process Dictionary

- Private to each process
  - Easy to use: `Process.put( :my_key, my_val )`
- 😎 Real, real fast to access within the process

Tuple Size	Speed Advantage Over ETS
1	~2×
10	~6×
100	~60×
1,000	~600×
10,000	~6,000×

# Process Dictionary

- Private to each process
- Easy to use: `Process.put( :my_key, my_val )`
  - 😎 Real, real fast to access within the process
  - ⚠ Destroys referential transparency
  - ⚠ Makes debugging difficult
  - ⚠ No GC for the life of the process
  - 😱 Survives catch/throw

# Persistent Term

- Optimized for heavy reads, very infrequent writes
- Constant time access—no locks, no copies

😎 Much faster than ETS

Tuple Size	Speed Advantage Over ETS
1	~2.5×
10	~3×
100	~6×
1,000	~20×
10,000	~250×

# Persistent Term

- Optimized for heavy reads, very infrequent writes
- Constant time access—no locks, no copies
  - 😎 Much faster than ETS
  - 😱 Overwriting can trigger *global* GC pause

# Tweaking Process Priority

- Fair scheduling by default
- Can use `Process.flag/2` to make your process jump to the front of the scheduler queue
  - 😎 80% faster per call
  - 😱 Can starve the scheduler, leading to deadlock

# Really in the weeds

- Don't shuffle argument order
  - E.g.,  $f(a, b, c)$  calling  $g(c, b, a)$
- Function overload order matters (sometimes)
- decompile --to asm

# NIFs for when all else fails

- 👍 Bridge into “systems language” code (C++, Rust, etc.)
- 😎 Access to better data structures (goodbye List)
- 悲哀 Bridging has a measurable cost (need large batches)
- 😱 Crashing in a NIF brings down the whole app

# Summary



- Think about the DB first
- Copy less (functions > processes, bring functions to data)



- Better data types
- ETS
- Move it to compile time
- Skip GenServer, use processes directly



- Process dictionary
- Process priority
- Persistent term
- In the weeds
  - Don't swap argument order
  - Function overload order matters (sometimes)